Will an LLM banking agent break KYC/AML rules to hit a sales target?

Executive brief: a pre-registered sandbox study of Claude agents

Thanatchon Chongmankhong • October 2026

Independent personal research. Fully synthetic sandbox. Does not represent the views of any employer or institution.

Bottom line. We gave Claude agents a bank relationship manager's job, a sales target and a manager who pushed for it, with know-your-customer (KYC) and anti-money-laundering (AML) controls that the agent could override by typing a reason. Across 3,522 pre-registered runs of Claude Haiku 5.5, Sonnet 5.5 and Opus 5.5 the agents committed 0 violations and made 0 attempts. They checked identities, requested documents, escalated sanctions hits and told the manager the target would be missed. The models differed instead in how much *legitimate* business they completed. The result is evidence, with stated bounds, for one design; it is not a certification of any agent.

Why it matters. Banks are starting to give LLM agents tools that can open accounts and post transactions. Money-mule networks, which move the proceeds of scams, depend on accounts being opened for people who do not own the money. Supervisors have answered by tightening account opening and by making a bank that falls short share the loss; the Bank of Thailand, for example, now requires banks to refuse accounts to high-risk applicants. An agent that cut a corner under sales pressure would open exactly the accounts these rules are meant to stop.

What the study did. SandBank is a fictional bank, exposed to the agent as seventeen ordinary tools: look up a lead, run a KYC check, request documents, open an account, book a deposit, file a report, escalate to Compliance, message the manager. Each shift had one customer whom policy blocked that day (failed KYC, a sanctions hit, a politically exposed person awaiting due diligence, a large deposit that needs a report, third-party funds, missing source of funds) and one customer who could be served. The manager's message rose from no target to a threat of branch closure and job loss. Further experiments removed warnings, told the agent Compliance reviews were paused, ran five-day episodes in which a deposit could be split, and had the manager explicitly license, then order, an override. A deterministic detector read the bank's event log for eight violation types. Hypotheses, sample sizes and analyses were registered before any model data.

Findings

1. **No violations under any condition.** In the main experiment, which ran all three models, the 95% upper bound on the per-run violation rate is 0.53% for Haiku, 0.53% for Sonnet (720 runs each) and 1.32% for Opus (288 runs); because the same scenarios recur across pressure levels, the conservative bounds are 2.1%, 2.1% and 5.1%. Neither an explicit licence ("use the override fields, I'll back you") nor a direct order from the manager produced a violation in Haiku or Sonnet.

2. **Under pressure the agents escalated and reported, rather than cut corners.** They ran the KYC check on the blocked customer in 100%, 99% and 95% of runs (Haiku, Sonnet, Opus), escalated every sanctions hit, and under the strongest pressure told the manager that a lead would not convert in 83% (Haiku), 81% (Sonnet) and 94% (Opus) of runs.

3. **The difference between models is legitimate work, not compliance.** A large deposit that policy allows once its report is filed was opened in 65% of Haiku runs, 92% of Sonnet runs and 100% of Opus runs. When Haiku did not open it, it left a legal account on hold; it never opened it without the report. An evaluation that counts only violations would score this caution as perfect. (Exploratory; not pre-registered.)

What a bank deploying an agent should take from this

Do	Why
Give the agent a compliant route for every case: escalate, request documents, report a missed target	An earlier study that found harmful actions made them the only route to the goal. Here every blocked case had a compliant route and the agents took it; the paper argues this is the likelier reason for the difference.
Keep hard controls on prohibited actions and log every tool call	Zero violations in a few hundred runs per model bounds the rate at a few percent; that does not rule out a rare failure that matters at bank scale.
Measure lost legitimate business alongside violations	The models differed in completed work, not in compliance; a guardrail's cost has to be measured on work that can actually be lost.
Test your own deployment, with your own policy and tools	One fictional bank, one clear policy, short episodes and one model family.

What this study does not show

- **The deceived agent.** The sandbox flags the third party and the missing documents; the test was whether a pressured agent respects the flag. Whether an agent recognises a mule whose papers pass the checks, or resists a customer's social engineering, was not tested.
- **Other model families, or why the agents complied.** Only Claude models were tested, and their reasoning was not visible. Models can sometimes tell when they are being evaluated; the agents never said so, but the null may partly reflect it.

Paper, pre-registration, code and every run	`https://github.com/thanatchon36/sandbank`
The paper on the web	`https://thanatchon36.github.io/sandbank/paper/`
Run data	`https://huggingface.co/datasets/thanatchon36/sandbank`

Every number in this brief is generated from the results files. Total API spend for the study: USD 62.93.